

Lab 5 CSC 453 Winter 2020 Can be done as a group of three students maximum

Program: Support for Lightweight Processes (liblwp.so)

This assignment requires you to implement support for lightweight processes (threads) under linux using the GNU C Compiler(gcc). A lightweight process is an independent thread of control— sequence of executed instructions—executing in the same address space as other lightweight processes. Here you will implement a non-preemptive user-level thread package.

This comes down to writing nine functions, described briefly , and in more detail below.

`lwp create(function,argument,stacksize)` create a new LWP
`lwp gettid()` return thread ID of the calling LWP
`lwp exit()` terminates the calling LWP
`lwp yield()` yield the CPU to another LWP
`lwp start()` start the LWP system
`lwp stop()` stop the LWP system
`lwp set scheduler(scheduler)` install a new scheduling function
`lwp get scheduler(void)` find out what the current scheduler is
`tid2thread(tid)` map a thread ID to a context

The Big Picture

Most of the magic will be in `lwp create()`. The job of `lwp create()` is to set up a thread's context so that when it is selected by the scheduler to run and one of `lwp yield()`, `lwp start()`, `lwp exit()` returns to it, it will start executing where you want it to. With a few minor differences, you'll find that `lwp yield()`, `lwp start()`, `lwp stop()`, and `lwp exit()` all more or less do the same thing: save one context, pick a thread to run, and load that thread's context.

Things to know

Everything in the rest of this document is intended to provide information needed to implement a lightweight processing package for a 64-bit Intel x86 64 CPU compiling with gcc.

Context: What defines a thread

Before we build a thread support library, we need to consider what defines a thread. Threads exist in the same memory as each other, so they can share their code and data segments, but each thread needs its own registers and stack to hold local data, function parameters, and return addresses.

Registers The x86 64 CPU (doing only integer arithmetic²) has sixteen registers of interest, shown in below:

`rax` General Purpose A `r8` General Purpose 8
`rbx` General Purpose B `r9` General Purpose 9
`rcx` General Purpose C `r10` General Purpose 10
`rdx` General Purpose D `r11` General Purpose 11
`rsi` Source Index `r12` General Purpose 12
`rdi` Destination Index `r13` General Purpose 13
`rbp` Base Pointer `r14` General Purpose 14
`rsp` Stack Pointer `r15` General Purpose 15

Since C has no way of naming registers, I have provided some useful tools below that will allow you to access these registers. First there are two macros, `Get SP()` and `Set SP()`, which will allow you to get or set the `%rsp` register directly. Second, is an assembly language file, `magic64.S` which contains a function `void swap rfiles(rfile *old, rfile *new)`. This does two things:

1. if `old != NULL` it saves the current values of all 16 registers to the struct registers pointed to by `old`.
2. if `new != NULL` it loads the 16 register values contained in the struct registers pointed to by `new` into the registers. For convenience, there are also two macros, `load context(c)` and `save context(c)` that do half of a swap should you find that desirable.

To assemble `magic64.S`, use gcc:

```
gcc -o magic64.o -c magic64.S
```

Floating Point State

As we said above, in addition to the registers, `swap rfiles()` also preserves the state of the x87 Floating Point Unit(FPU). This is stored in the last element of the struct `rfile`, the struct `fxsave` called `fxsave`. This structure holds all the FPU state. Important: when you initialize your thread's register file, you will have to initialize this structure to the predefined value `FPU_INIT` like so:

```
newthread->state.fxsave=FPU_INIT;
```

Stack structure: The gcc calling convention

The extra registers available to the x86 64 allow it to pass some parameters in registers. This makes the overall calling convention a little more complicated, but, in practice, it will be easier for your program since you won't be passing enough parameters to push you out of the registers onto the stack.

The steps of the convention are as follows (illustrated in Figures 1a–f):

a. Before the call Caller places the first six integer arguments into registers `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, and `%r9`. If there are more, they are pushed onto the stack in reverse order. This is shown in the figure, but you won't encounter more in this assignment.

b. After the call The call instruction has pushed the return address onto the stack.

c. Before the function body If the function has more parameters and local variables than will fit into the registers it will execute the following two instructions to set up its frame:

```
pushq %rbp
```

```
movq %rsp,%rbp
```

Then, it adjusts the stack pointer to leave room for any locals it may need.

d. Before the return Before returning, the function needs to clean up after itself. To do this, before returning it executes a leave instruction. This instruction is equivalent to:

```
movq %rbp,%rsp
```

```
popq %rbp
```

The effect is to rewind the stack back to its state right after the call.

e. After the return After the return, the Return address has been popped off the stack, leaving it looking just like it did before the call.

Remember, the `ret` instruction, while called "return", really means "pop the top of the stack into the program counter."

f. After the cleanup Finally, the caller pops off any parameters on the stack and leaves the stack is just like it was before.

LWP system architecture

Everything you need is defined in `lwp.h`, `fp.h` and `magic64.S`, included in Figures 2 and 3. At the heart of `lwp.h` is the definition of a struct `threadinfo_t` which defines a thread's context. This contains:

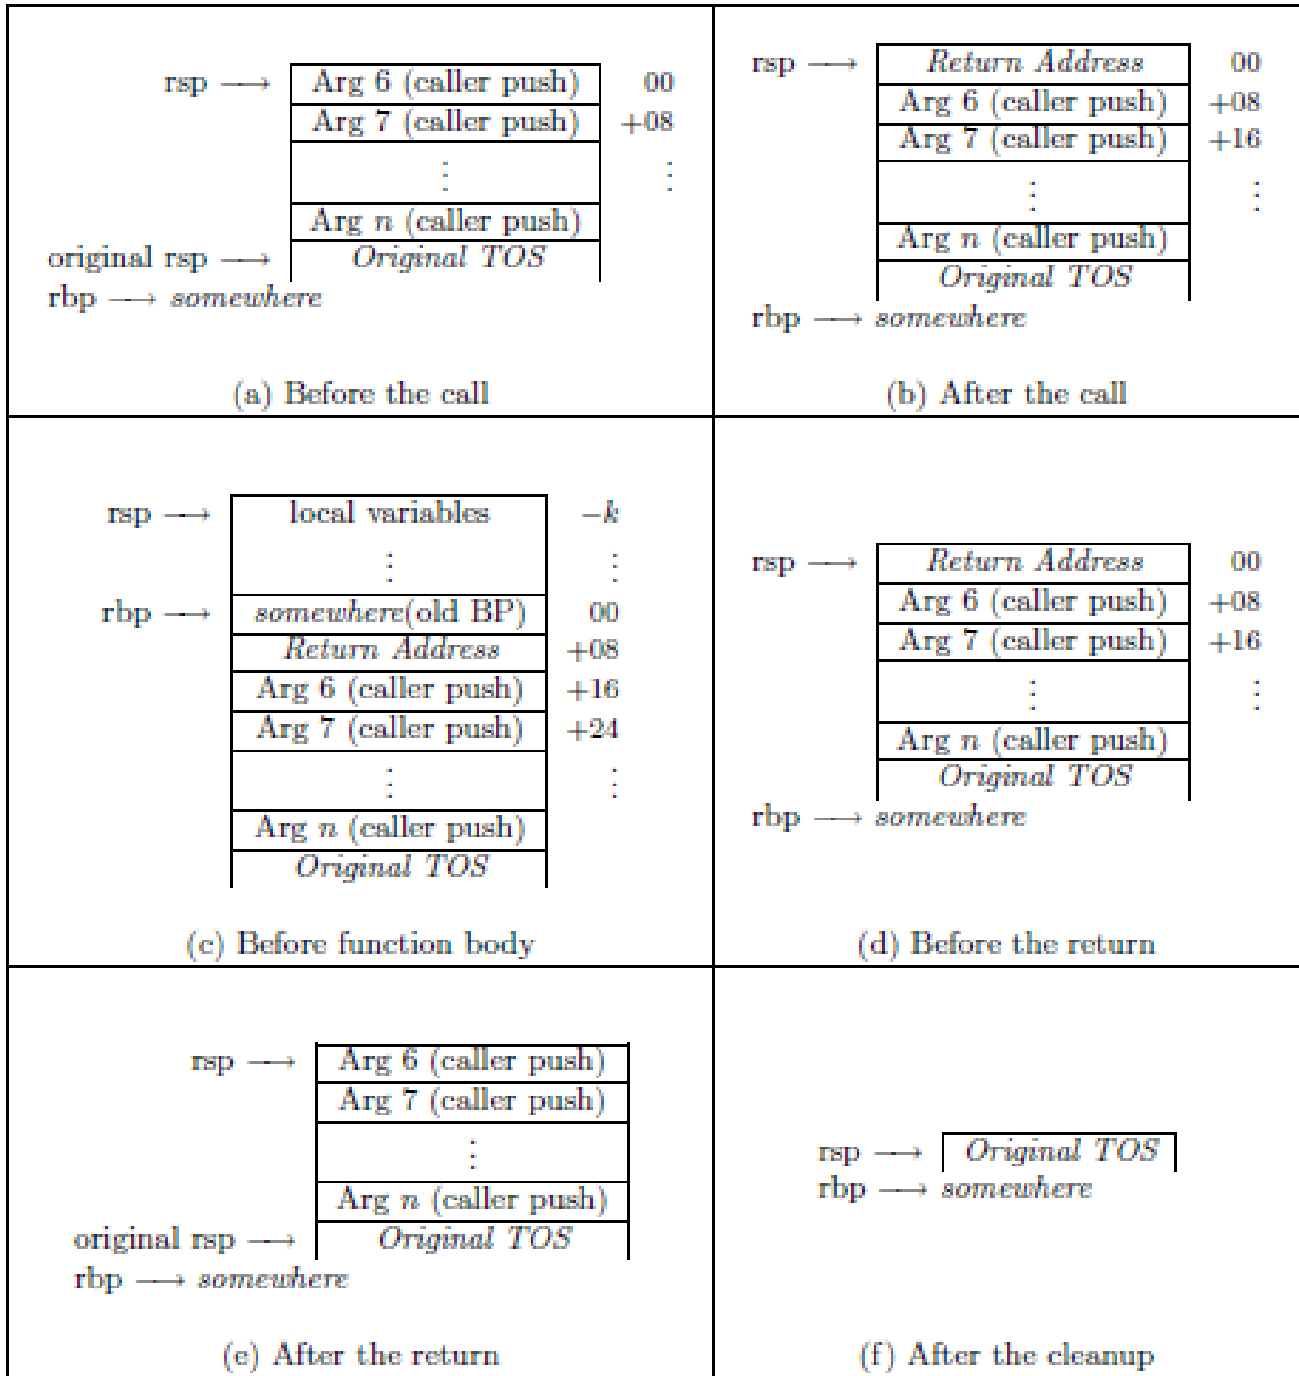


Figure 1: Stack development (Remember that the real stack is upside-down)

```

#ifndef LWP_H
#define LWP_H
#include <sys/types.h>

#ifndef TRUE
#define TRUE 1
#endif
#ifndef FALSE
#define FALSE 0
#endif

#if defined(_x86_64)
#include "fp.h"
typedef struct __attribute__((aligned(16))) __attribute__((packed))
registers {
    unsigned long rax; /* the sixteen architecturally-visible regs. */
    unsigned long rbx;
    unsigned long rcx;
    unsigned long rdx;
    unsigned long rsi;
    unsigned long rdi;
    unsigned long rbp;
    unsigned long rsp;
    unsigned long r8;
    unsigned long r9;
    unsigned long r10;
    unsigned long r11;
    unsigned long r12;
    unsigned long r13;
    unsigned long r14;
    unsigned long r15;
    struct fsave fsave; /* space to save floating point state */
} rfile;
#else
#error "This only works on x86_64 for now"
#endif

typedef unsigned long tid_t;
#define NO_THREAD 0 /* an always invalid thread id */

typedef struct threadinfo st *thread;
typedef struct threadinfo st {
    tid_t tid; /* lightweight process id */
    unsigned long *stack; /* Base of allocated stack */
    size_t stacksize; /* Size of allocated stack */
    rfile rfile; /* saved registers */
    thread lib_one; /* Two pointers reserved */
    thread lib_two; /* for use by the library */
    thread sched_one; /* Two more for */
    thread sched_two; /* schedulers to use */
} context;

typedef void (*lwpfun)(void *); /* type for lwp function */

/* Tuple that describes a scheduler */
typedef struct scheduler {
    void (*init)(void); /* initialize any structures */
    void (*shutdown)(void); /* tear down any structures */
    void (*admit)(thread new); /* add a thread to the pool */
    void (*remove)(thread victim); /* remove a thread from the pool */
    thread (*next)(); /* select a thread to schedule */
} *scheduler;

/* lwp functions */
extern tid_t lwp_create(lwpfun, void *, size_t);
extern void lwp_exit(void);
extern tid_t lwp_gettid(void);
extern void lwp_yield(void);
extern void lwp_start(void);
extern void lwp_stop(void);
extern void lwp_set_scheduler(scheduler fun);
extern scheduler lwp_get_scheduler(void);
extern thread tid2thread(tid_t tid);

/* Macros for stack pointer manipulation:
 *
 * GetSP(var) Sets the given variable to the current value of the
 * stack pointer.
 * SetSP(var) Sets the stack pointer to the current value of the
 * given variable.
 */
#if defined(_x86_64) /* X86 only code */
#define BAIL_SIGNAL SIGSTKFLT
#define GetSP(sp) asm("movq %rsp, %0": "=r" (sp) : )
#define SetSP(sp) asm("movq %0, %rsp": : "r" (sp) : )
#else /* END x86 only code */
#error "This stack manipulation code can only be compiled on an x86"
#endif

#if defined(_APPLE_)
#undef BAIL_SIGNAL
#define BAIL_SIGNAL SIGABRT
#endif

/* prototypes for asm functions */
#define load_context(c) (swap_rfiles(NULL, c))
#define save_context(c) (swap_rfiles(c, NULL))
void swap_rfiles(rfile *, rfile *to);

```

Figure 2: Definitions and prototypes for LWP: lwp.h

The thread's thread ID. This must be a unique integer that stays the same for the lifetime of the thread. It's what a thread may use to identify itself. (NO_THREAD is defined to be 0 and is always invalid.) You may assume that there will never be more than $2^{64} - 2$ threads, so a counter is just fine.

- A pointer to the base of the thread's allocated stack space so that it can later be free()d.
 - A struct registers that contains a copy of all the thread's stored registers.
 - Four pointers:
 - lib one and lib two are reserved for the use of the library internally. (E.g., to maintain a global linked list of all valid threads for implementing tid2thread().)
 - sched one and sched two are reserved for use by schedulers.
- Together, these hold all the state we need for each thread.

Scheduling

The lwp scheduler is a structure that holds pointers to five functions. These are:

- void init(void) This is to be called before any threads are admitted to the scheduler. It's to allow the scheduler to set up. This one is allowed, to be NULL, so don't call it if it is.
 - void shutdown(void) This is to be called when the lwp library is done with a scheduler to allow it to clean up. This, too, is allowed, to be NULL, so don't call it if it is.
 - void admit(thread new) Add the passed context to the scheduler's scheduling pool.
 - void remove(thread victim) Remove the passed context from the scheduler's scheduling pool.
 - thread next() Return the next thread to be run or NULL if there isn't one.
- Changing schedulers will involve initializing the new one, pulling out all the threads from the old one (using next() and remove()) and admitting them to the new one (with admit()), then shutting down the old scheduler.

A note on function pointers:

Remember, the name of a function is its address, so you can pass a pointer to a function just by using its name. For example, my round robin scheduler is defined like so:

```
static struct scheduler rr publish = {NULL, NULL, rr admit, rr remove, rr next};
scheduler RoundRobin = &rr publish;
Calling a function pointer is just a matter of dereferencing it and applying it to an argument.
E.g.:
thread nxt;
nxt = RoundRobin->next()
```

How to get started

1. Allocate a stack for each LWP.
 2. Build a stack frame on the stack so that when that context is loaded it will properly return to the lwp's function with the stack and registers arranged as it will expect. This involves making the stack look as if the thread called you and was suspended.
 3. When lwp start() is called:
 - (a) save the "real" context somewhere where lwp stop() can find it,
 - (b) pick one of the lightweight processes to run (by calling the scheduler).
 - (c) Load its context with swap rfiles() and you should be off and running.
- Remember, what you are trying to do is to build a context so that when lwp yield() selects it, loads its registers, and returns, it starts executing its very first instruction with the stack pointer pointing to a stack that looks like it had just been called.
- If the arguments fit into registers, this will simply be:

```
rsp → 

|                |     |
|----------------|-----|
| Return Address | 00  |
| Original TOS   | +08 |


rbp → somewhere
```

But what is this return address? It's supposed to be the place where the thread function should go "back" to after it's done, but it didn't come from anywhere. Use lwp exit(). That way either it calls lwp exit() or it returns there, but one way or the other when it's done, lwp exit() will be called.

Tricks, Tools, and Useful Notes

- a segmentation violation may mean
 - a stack overflow.
 - stack corruption
 - all the other usual causes
- Use the CSL linux machines
- But I really want to use my Mac.

Ok. . . but there are a few things that are different about doing this under OSX:

- OSX requires all stack frames to be 16-byte aligned.
- Dynamic libraries have the suffix .dylib
- The path the loader searches for dynamic libraries is DYLD_LIBRARY_PATH.
- It is possible to compile multiple architectures of library into a single .dylib file. See lido(1) for details.
- Finally, you'll need to be sure it compiles and runs on Linux, since that's where it'll be graded.

<code>lwp_create(function,argument,stacksize);</code>	Creates a new lightweight process which executes the given function with the given argument. The new processes's stack will be <code>stacksize</code> words. <code>lwp_create()</code> returns the (lightweight) thread id of the new thread or <code>-1</code> if the thread cannot be created.
<code>lwp_gettid(void);</code>	Returns the tid of the calling LWP or <code>NO_THREAD</code> if not called by a LWP.
<code>lwp_yield(void);</code>	Yields control to another LWP. Which one depends on the scheduler. Saves the current LWP's context, picks the next one, restores that thread's context, and returns.
<code>lwp_exit(void);</code>	Terminates the current LWP and frees its resources. Calls <code>sched->next()</code> to get the next thread. If there are no other threads, calls <code>lwp_stop()</code> .
<code>lwp_start(void);</code>	Starts the LWP system. Saves the original context (for <code>lwp_stop()</code> to use later), picks a LWP and starts it running. If there are no LWPs, returns immediately.
<code>lwp_stop(void);</code>	Stops the LWP system, restores the original stack pointer and returns to that context. (Wherever <code>lwp_start()</code> was called from. <code>lwp_stop()</code> does not destroy any existing contexts, and thread processing will be restarted by a call to <code>lwp_start()</code>).
<code>lwp_set_scheduler(scheduler);</code>	Causes the LWP package to use the given <code>scheduler</code> to choose the next process to run. Transfers all threads from the old scheduler to the new one in <code>next()</code> order. If <code>scheduler</code> is <code>NULL</code> , or has never been set, the scheduler should do round-robin scheduling.
<code>lwp_get_scheduler(void);</code>	Returns the pointer to the current scheduler.
<code>tid2thread(tid);</code>	Returns the thread corresponding to the given thread ID, or <code>NULL</code> if the ID is invalid

Table 3: The LWP functions

If you want to find out what your compiler is really doing, use the `gcc -S` switch to dump the assembly output.

`gcc -S foo.c`

will produce `foo.s` containing all the assembly.

- `lwp_exit()` is probably the second trickiest part of this assignment because you must be careful not to `free()` a stack that you're still using.
- Remember that stacks start in high memory and grow towards low memory.
- Also remember that pointer arithmetic is done in terms of the size of the thing pointed-to.
- Instructions for building and using shared libraries are included in Asgn1 if you need to review.
- Note that `lwp_stop()` does not necessarily mean stop forever. It should be possible to call `lwp_stop()` then later call `lwp_start()` to restart.
- Finally, remember that there doesn't have to be a next thread. If `sched->next()` returns `NULL`, `lwp_yield()`, `lwp_exit()`, and `lwp_start()` should restore the original system context. (as `lwp_stop()` does).

Supplied Code

There are several pieces of supplied code

File Description/Location

`lwp.h` Header file for `lwp.c`

fp.h Header file for preserving floating point state
libPLN.a precompiled library of lwp functions (for testing)
libsnakes.a precompiled library of snake functions
magic64.S ASM source for swap rfiles()
snakes.h header file for snake functions
hungrymain.c demo program for hungry snakes
snakemain.c demo program for wandering snakes
numbersmain.c demo program with indented numbers

Note: When linking with libsnakes.a it is also necessary to link with the standard library ncurses using -lncurses on the link line. Ncurses is a library that supports text terminal manipulation.

What to turn in

Submit to polylearn

- your well-documented source file(s).
 - Your header file, lwp.h, suitable for inclusion with other programs. This must be compatible with the distributed one, but you may extend it. A makefile (called Makefile) that will build liblwp.so from your source when invoked with no target or with the target "liblwp.so".
 - A README file that contains:
 - Your name(s), including your login name(s) in parentheses (e.g. "(pnico)").
 - Any special instructions.
 - Any other thing you want me to know while I am grading it.
- The README file should be plain text, i.e, not a Word document, and should be named "README", all capitals with no extension.

Sample runs

We did these in class. If you want, though, you can use the provided libPLN.a to build your own samples.